

Process Management

- An executable program for a CPU architecture consists of a sequence of instructions acceptable to the CPU. Executable **programs reside on a media**.
- A **Process** is an executable **program in execution**.
 - ▶ Program **loaded into memory** area called process address space.
 - ▶ **Address space -- Set of memory locations accessible to the loaded program**. Typically, a process has its own address space.
 - ▶ CPU fetches instructions from the address space to execute them.
 - ▶ The data part of the process is also stored in the address space.
 - ▶ A **pointer to the next instruction** to be executed is maintained.
- OS terminology for a process - task.
- UNIX Operating Systems are multi-tasking:
 - ▶ Multiple programs are loaded into memory and CPU is switched among them.

Process Management

- UNIX Operating Systems are multi-user:
 - ▶ Programs of multiple users are loaded in memory.
- UNIX Operating Systems are time sharing:
 - ▶ CPU switching is done so that user interactive programs generate quick responses to the users.
- Multiple processes can be created simultaneously from a single program.
 - ▶ Multiple vi processes, each created from /usr/bin/vi, can be executing concurrently on behalf of more than one user. The file /usr/bin/vi it self will be owned by root, but has rest of the world execute permission.
- Each process is one executing instance of the program from which it is created.

Process Identifiers

Identifier	System call to fetch ID
Process ID	<code>pid_t getpid(void)</code>
Parent process ID	<code>pid_t getppid(void)</code>
Real user ID	<code>uid_t getuid(void);</code>
Effective user ID	<code>uid_t geteuid(void);</code>
Real group ID	<code>gid_t getgid(void);</code>
Effective group ID	<code>gid_t getegid(void);</code>
Process group ID	<code>pid_t getpgrp(void);</code>
Supplementary group IDs	<code>int getgroups()</code>

Process Creation


- A process can be created (only by) another process.
- First user process is hand crafted by kernel after it is loaded by boot strapping.
- Referred to as init process and has $PID = 1$.
- The init process is responsible for initializing (bringing up) the system.
- The lineage of every user process can be traced to init.

Process Creation

`pid_t fork(void);`

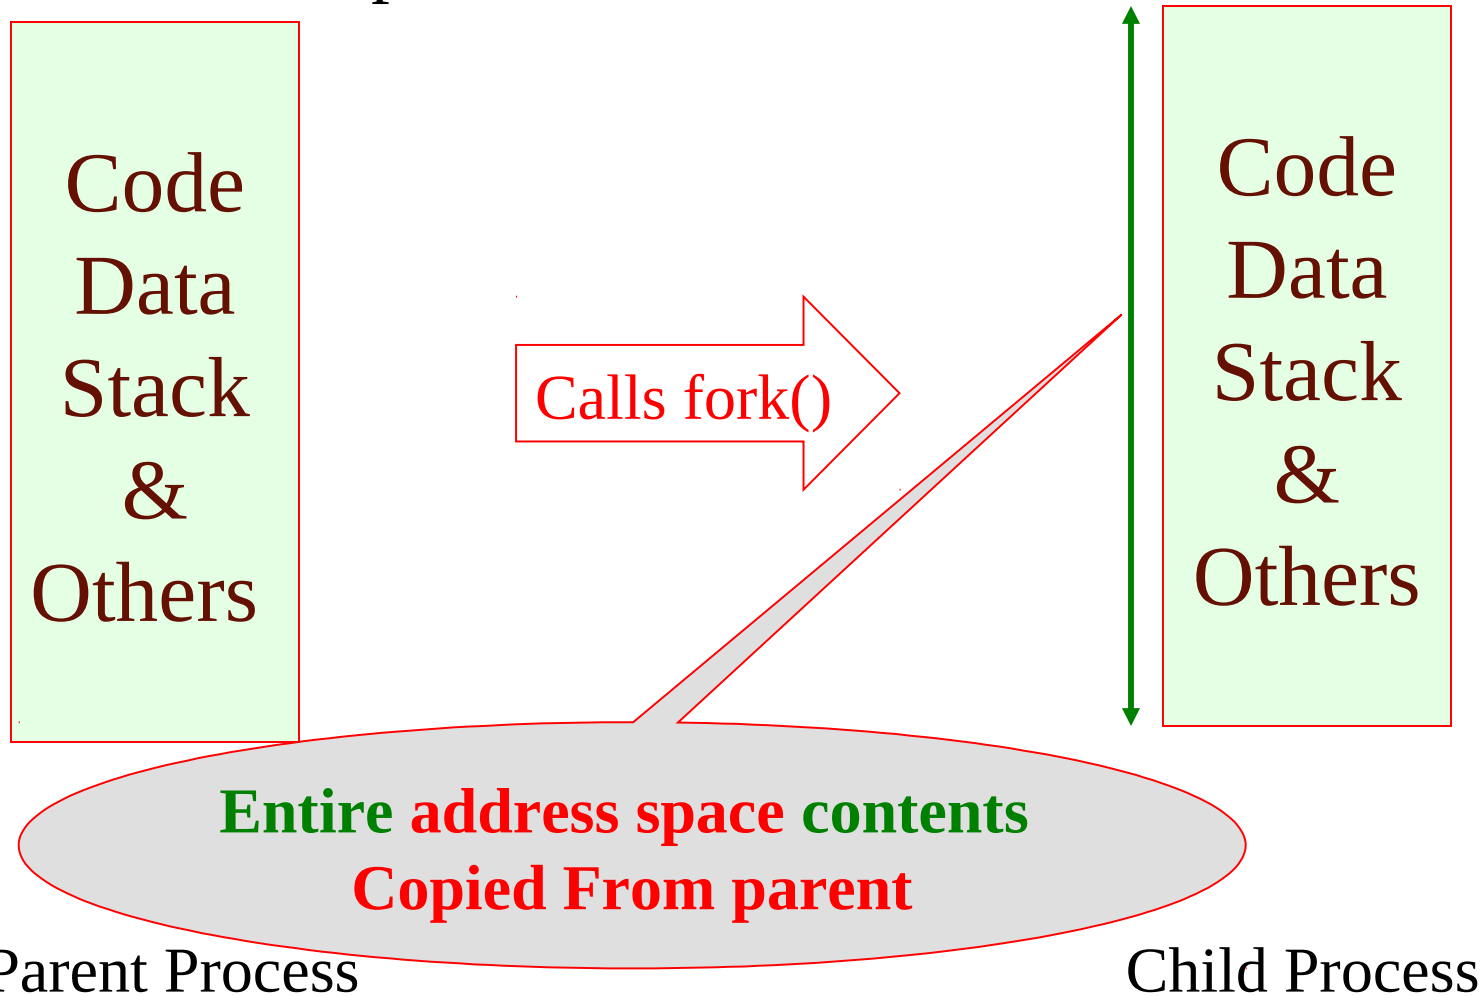
- Will create a new child process with the calling process being the parent.
- By the time `fork()` returns, both parent and child process will be eligible to run.

A CPU can run only one process at a time.



Process Creation

- Parent process's address space will be copied to the address space of the child.



Process Creation

- Child's address space is an exact replica of parent's. This includes **code, data and stack**.
- Child Starts execution as if `fork()` has just returned.
- Return value of `fork()`:
 - In parent -- new Child's Process ID.
 - In Child -- Zero

```
short X= 100;
main(){
int Y=20;
X++, Y--;
pid_t fork_ret = fork();
printf( "pid= %d\n", fork_ret );
printf("X=%d,Y=%d\n", X, Y);
}
```

**Parent calls
fork()**

```
short X= 100;
main(){
int Y=20;
X++, Y--;
pid_t fork_ret = fork();
printf( "pid= %d\n", fork_ret );
printf("X=%d,Y=%d\n", X, Y);
}
```

**Child starts
at assignment.**

Different Execution Paths

```
main(){
pid_t fork_ret = fork();

if (fork_ret < 0)
fprintf(stderr, "Error forking\n");

if (fork_ret > 0) {
    printf( "Parent Process:\n");
    printf( "PID=%d\n", getpid());
}

if (fork_ret == 0) {
    printf( "Child Process:\n");
    printf( "PID=%d\n", getpid());
}
}
```

```
main(){
pid_t fork_ret = fork();

if (fork_ret < 0)
fprintf(stderr, "Error forking\n");

if (fork_ret > 0) {
    printf( "Parent Process:\n");
    printf( "PID=%d\n", getpid());
}

if (fork_ret == 0) {
    printf( "Child Process:\n");
    printf( "PID=%d\n", getpid());
}
}
```


Different Execution Paths

```
main(){
pid_t fork_ret = fork();

pid_t pid= getpid();
int Count = 40;

if (fork_ret > 0) {
    printf( "Parent pid=%d:\n", pid);
    printf( "Count=%d\n", ++Count);
}

if (fork_ret == 0) {
    printf( "Child pid:%d\n", pid);
    printf( "Count=%d\n",--Count);
}
}
```

```
main(){
pid_t fork_ret = fork();

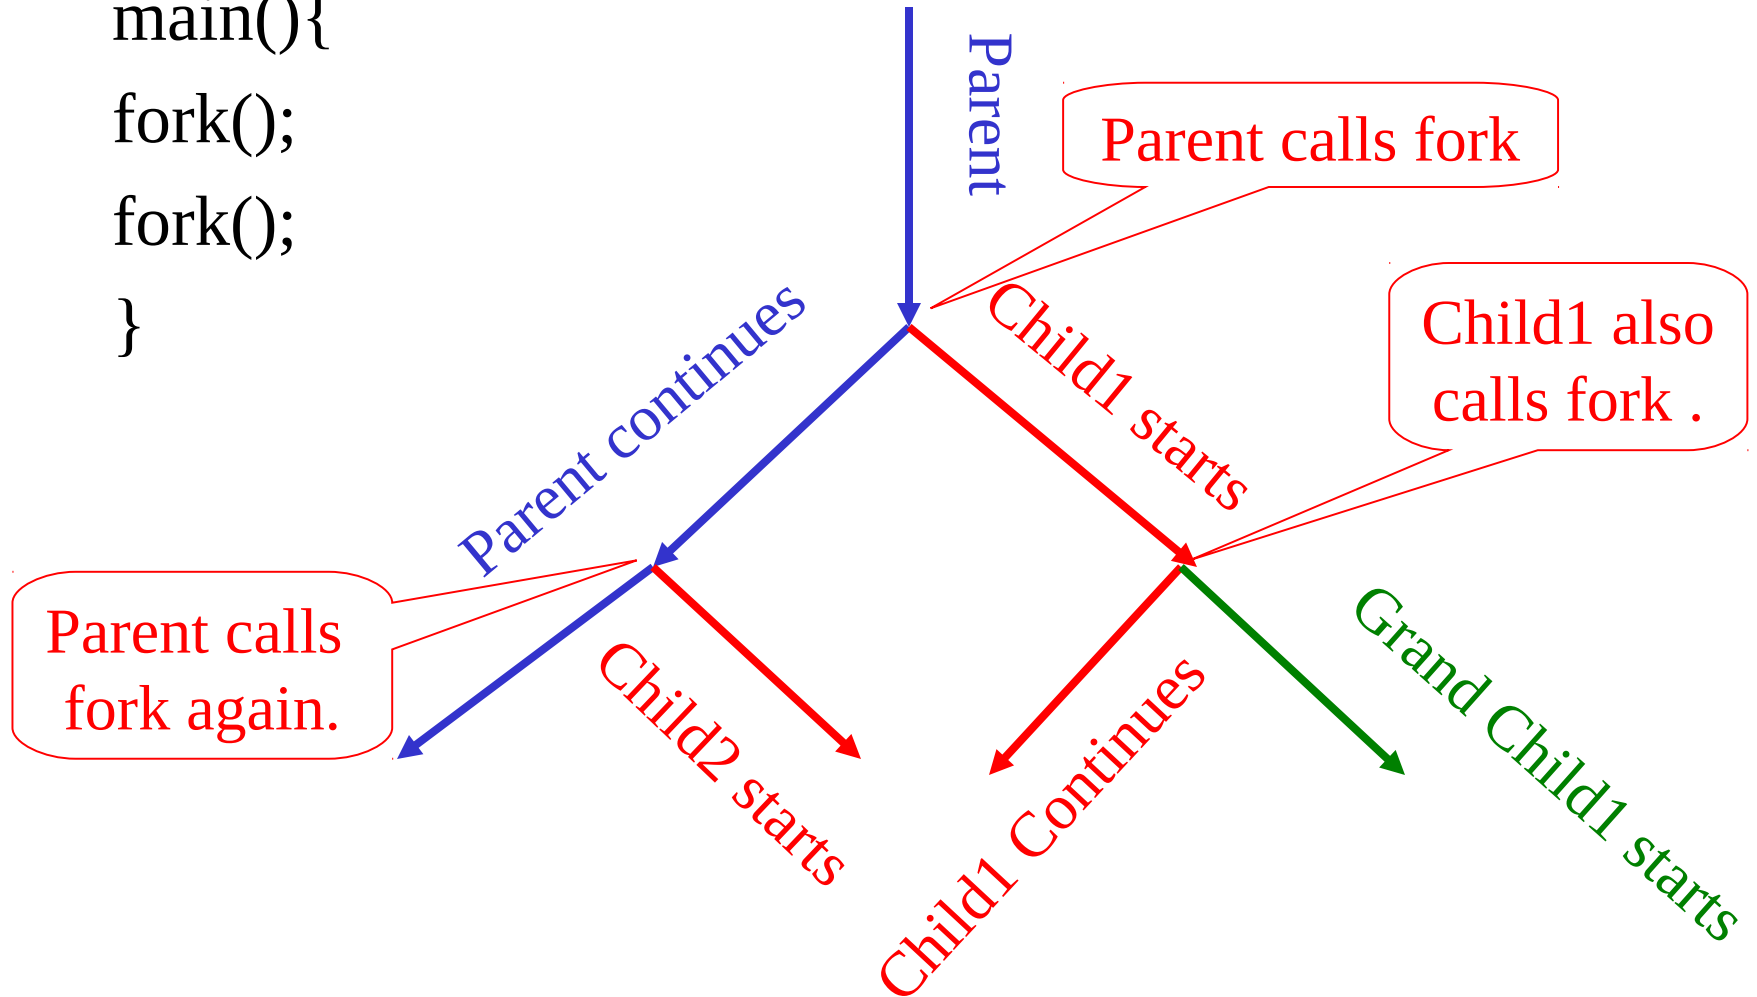
pid_t pid= getpid();
int Count = 40;

if (fork_ret > 0) {
    printf( "Parent pid=%d:\n", pid);
    printf( "Count=%d\n",++Count);
}

if (fork_ret == 0) {
    printf( "Child pid:%d\n", pid);
    printf( "Count=%d\n",--Count);
}
}
```

How many Processes created?

```
main(){  
fork();  
fork();  
}
```



FDT duplication on fork

```
main(){
char buf[256];
int fd1=open("f1",O_RDWR);
pid_t fork_ret = fork();

if (fork_ret > 0) { /* Parent */
read(fd1, buf, 80);
close(fd1);
int fd2=open("f2",O_RDONLY);
    /* Only FDT Entry Freed.*/

if (fork_ret == 0) { /* Child */
write(fd1, buf, 60);
int fd3=open("f3",O_WRONLY);
    }
}
```

FDT duplication on fork

FDT of Parent before fork

0	FD Flags, Pointer to FTE of stdin device.
1	FD Flags, Pointer to FTE of stdout device.
2	FD Flags, Pointer to FTE of stderr device.
3	FD Flags, Pointer to FTE of file “f1”

FDT of child on creation - Duplicated from Parent

0	FD Flags, Pointer to FTE of stdin device.
1	FD Flags, Pointer to FTE of stdout device.
2	FD Flags, Pointer to FTE of stderr device.
3	FD Flags, Pointer to FTE of file “f1”

FDT duplication on fork

```
main(){  
char buf[256];  
int fd1=open("f1",O_RDWR);  
pid_t fork_ret = fork();
```

Parent's FDT

Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE

Child's FDT

Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE

Global File Table

Unused
Unused
Unused
O_RDWR, 0
Unused
Unused
Unused
Unused

```
}
```

11/11/14

Process Management

Logic Option

FDT duplication on fork

```
main(){  
    char buf[256];  
    int fd1=open("f1",O_RDWR);  
    pid_t fork_ret = fork();  
  
    if (fork_ret > 0) { /* Parent */  
        read(fd1, buf, 80);  
    }  
}
```

Parent's FDT

Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE

Child's FDT

Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE

Global File Table

Unused
Unused
Unused
O_RDWR, 80
Unused
Unused
Unused
Unused

Logic Option

FDT duplication on fork

```
main(){
char buf[256];
int fd1=open("f1",O_RDWR);
pid_t fork_ret = fork();

if (fork_ret > 0) { /* Parent */
read(fd1, buf, 80);
}

if (fork_ret == 0) { /* Child */
write(fd1, buf, 60);
}
}
```

Parent's FDT

Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE

**No Guarantee that
parent executes first.**

Child's FDT

Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE

Global File Table

Unused
Unused
Unused
O_RDWR, 140
Unused
Unused
Unused
Unused

FDT duplication on fork

```
main(){
char buf[256];
int fd1=open("f1",O_RDWR);
pid_t fork_ret = fork();

if (fork_ret > 0) { /* Parent */
read(fd1, buf, 80);
close(fd1);
}/* Only FDT Entry Freed.*/

if (fork_ret == 0) { /* Child */
write(fd1, buf, 60);
}

}
```

Parent's FDT

Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE
FREE for Reuse

Child's FDT

Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE

Global File Table

Unused
Unused
Unused
O_RDWR, 140
Unused
Unused
Unused
Unused

Logic Option

FDT duplication on fork

```
main(){
char buf[256];
int fd1=open("f1",O_RDWR);
pid_t fork_ret = fork();

if (fork_ret > 0) { /* Parent */
read(fd1, buf, 80);
close(fd1);
int fd2=open("f2",O_RDONLY);
}          /* fd2 only in parent */
if (fork_ret == 0) { /* Child */
write(fd1, buf, 60);
}

}
```

Parent's FDT

Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE

Child's FDT

Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE

Global File Table

O_RDONLY,0
Unused
Unused
O_RDWR, 140
Unused
Unused
Unused
Unused

Logic Option

FDT duplication on fork

```
main(){
char buf[256];
int fd1=open("f1",O_RDWR);
pid_t fork_ret = fork();

if (fork_ret > 0) { /* Parent */
read(fd1, buf, 80);
close(fd1);
int fd2=open("f2",O_RDONLY);
}          /* fd2 only in parent */
if (fork_ret == 0) { /* Child */
write(fd1, buf, 60);
int fd3=open("f3",O_WRONLY);
}          /* fd3 only in Child */
}
```

Parent's FDT

Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE

Child's FDT

Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE

Global File Table

O_RDONLY,0
Unused
Unused
O_RDWR, 140
Unused
Unused
O_WRONLY,0
Unused



FDT duplication on fork

```
main(){  
char buf[256];  
int fd1=open("f1",O_RDWR);  
pid_t fork_ret = fork();  
  
int fd2=open("f2",O_RDONLY);  
  
}
```

**File
descriptor(s)
in which process?**

Parent's FDT

Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE

Child's FDT

Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE
Fdflags, ptr To FTE

Global File Table

O_RDONLY,0
Unused
Unused
O_RDWR,140
Unused
Unused
O_RDONLY,0
Unused

FDT: Parent and child Open

FDT of Parent After Opening f2.

0	FD Flags, Pointer to FTE of stdin device.
1	FD Flags, Pointer to FTE of stdout device.
2	FD Flags, Pointer to FTE of stderr device.
3	FD Flags, Pointer to FTE of file “f1”
4	FD Flags, Pointer to FTE of file “f2”

FDT of child After Opening f3.

0	FD Flags, Pointer to FTE of stdin device.
1	FD Flags, Pointer to FTE of stdout device.
2	FD Flags, Pointer to FTE of stderr device.
3	FD Flags, Pointer to FTE of file “f1”
4	FD Flags, Pointer to FTE of file “f3”

Process Termination

- Normal Termination
 - ▶ Calls `exit()`
 - ▶ Calls `_exit();`
 - ▶ Returns from `main()`
- Abnormal termination
 - ▶ **Process terminated due to a signal.**
SIGTERM, SIGUSR1, SIGKILL
 - ▶ **The job of the process is partially completed.**
- Process Suspension
 - ▶ **Process suspended by a signal.**
SIGSTOP, SIGTTIO
 - ▶ **Execution can be resumed by SIGCONT signal.**

Normal Termination

- `void _exit(int status);`
 - ▶ **System call that terminates the process normally.**
 - ▶ **Closes all open descriptors.**
 - ▶ **Status % 256 is the exit status of the process.**
- `void exit(int status);`
 - ▶ **Library call.**
 - ▶ **Flushes all open streams and closes them.**
 - ▶ **Status % 256 is the exit status.**
 - ▶ **Calls `_exit()` with LSB of status.**
- `return (status);` from `main()`.
 - ▶ **Equivalent to calling `exit(status);`**

Normal Termination ?

```
main() {  
    int fd = open ( "f1", O_RDWR);  
    if( fd < 0) {  
        fprintf(stderr, " Error opening f1\n");  
        _exit(10);  
    }  
}
```

- Normal termination -- Process terminates itself.
- Abnormal termination -- If the process is terminated by an external signal before it completes its execution.

Process Termination status

- When a process terminates, the kernel registers its termination status.
- Termination status is generally of integral type.
- Termination status embeds within it the process's termination type and the reason for termination.
 - ▶ Exits Status - Normal Termination
 - ▶ Signal Num - Abnormal Termination
 - ▶ Sigan Num - Process Suspension

Retrieving Termination Status

- `pid_t wait(int *status_ptr);`

**Only Parent
can retrieve.**


- Parent process can retrieve the termination status of a child with `wait()` system call.
- The terminated **child's PID** is returned by `wait`.
- The terminated **child's termination status** is stored in the location pointed at by `status_ptr`.

Termination Status - Embeds....



**Normal
Termination**

**Abnormal
Termination.**



**Implementation
Dependent usage
of byte positions**

**Process
Suspended**

Termination Status -- Study

- **WIFEXITED(*status_ptr)**
 - ▶ True if child process had normal termination.
 - ▶ **WEXITSTATUS(* status_ptr)** - Gives exit status.
- **WIFSIGNALED(* status_ptr)**
 - ▶ True if child process had abnormal termination.
 - ▶ **WTERMSIG(* status_ptr)** - Gives signal number.
- **WIFSTOPPED(* status_ptr)**
 - ▶ True if child process was suspended.
 - ▶ **WSTOPSIG(* status_ptr)** - Gives signal number.
 - ▶ Only with WUNTRACED option of waitpid().



Use these macros
for portability

Retrieving Termination Status

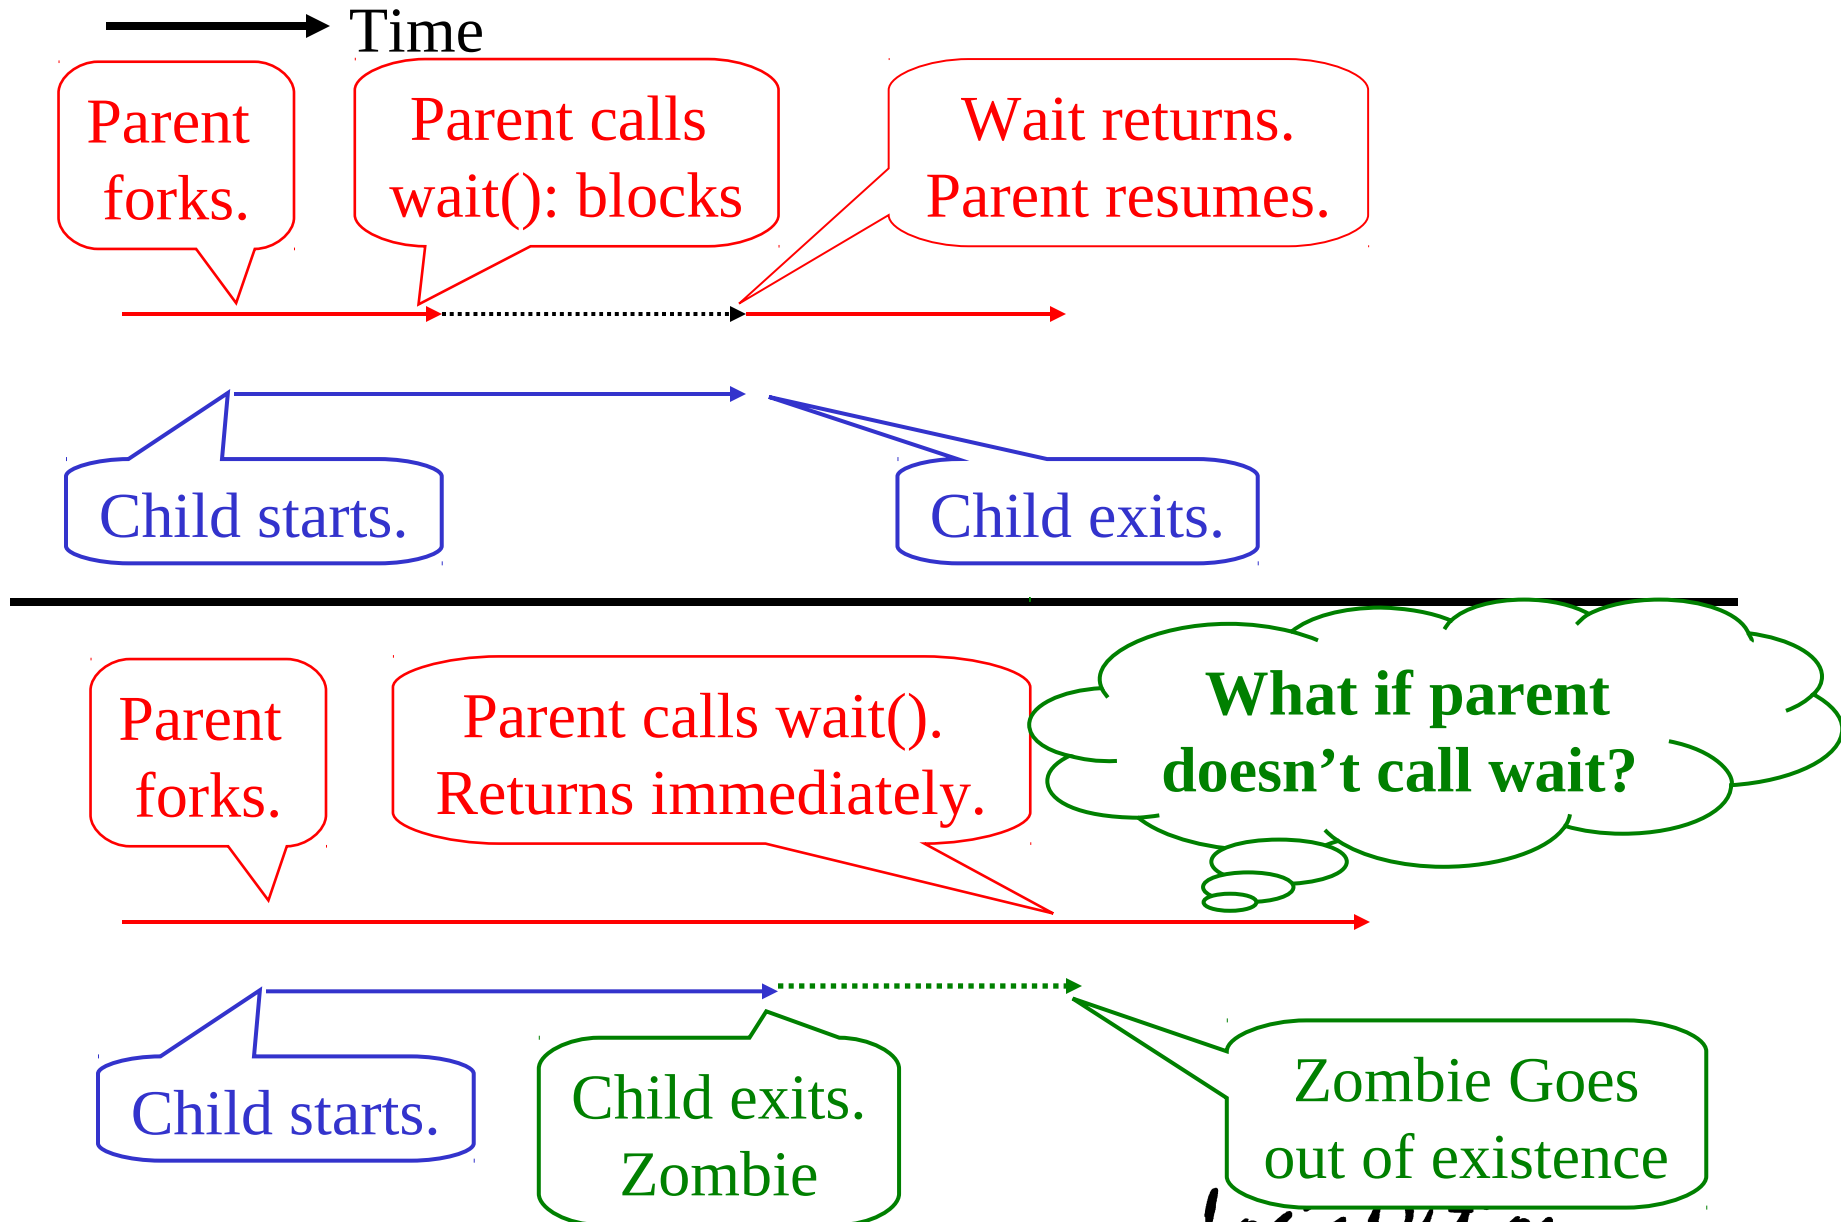
```
main(){  
  if(fork())  
  {  
    int s;  
    pid_t pid=wait(&s);  
    if(WIFEXITED(s))  
      printf("Child Process ID=%d,\n  
             Exit status=%d\n",  
             pid, WEXITSTATUS(s));  
  }  
  else{  
    _exit(10);  
  }  
}
```

Termination Status
of child with embedded
Exit status

KERNEL

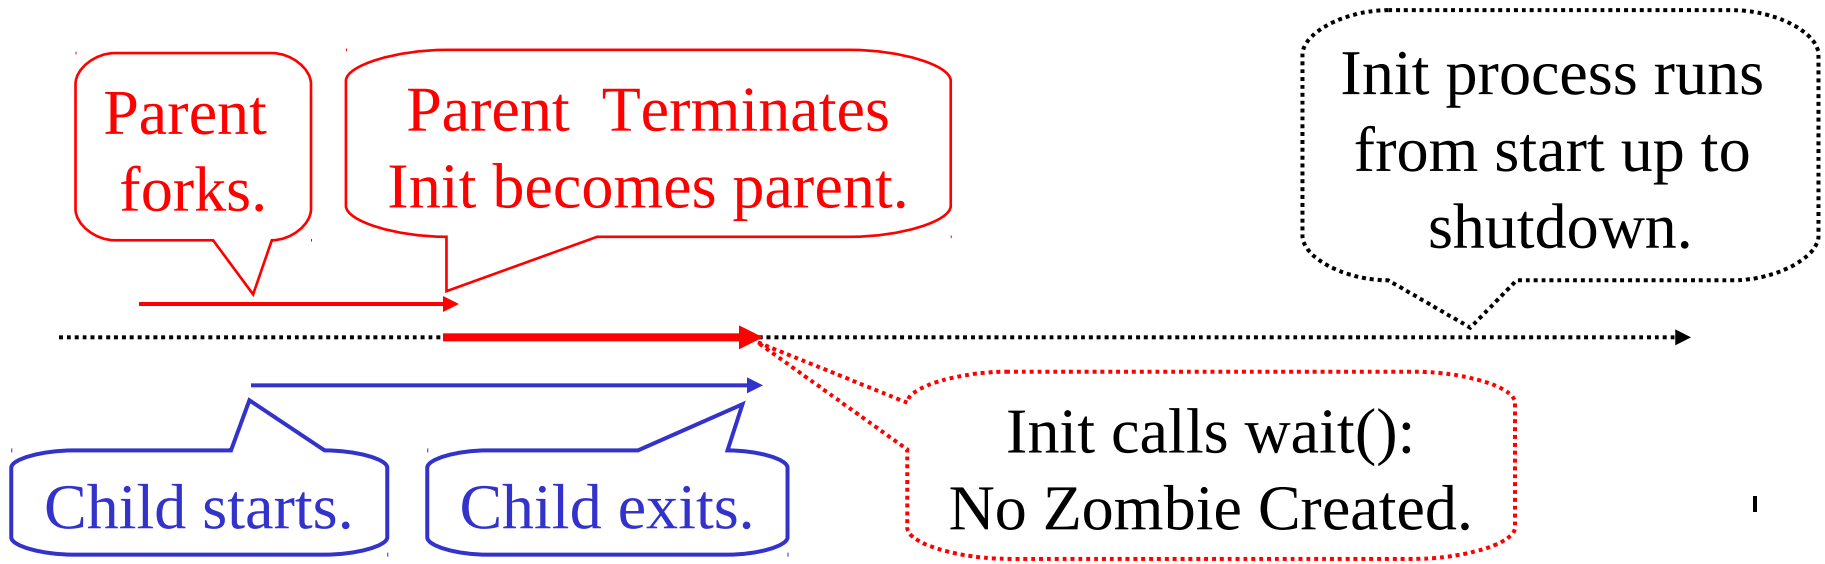
Child's
Exit Status

Child Terminates before Parent

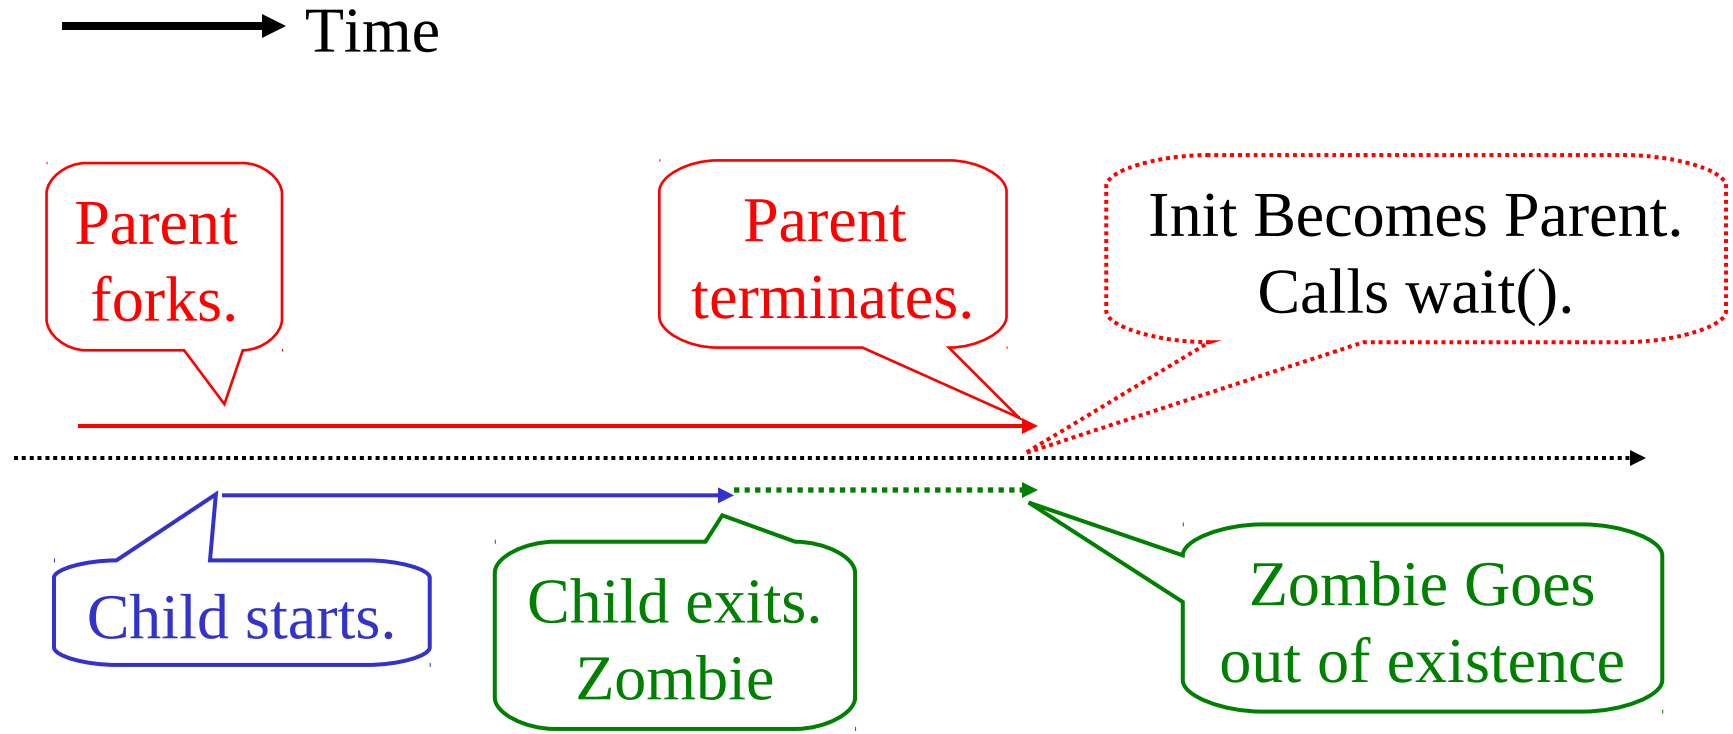


Parent Terminates before Child

→ Time

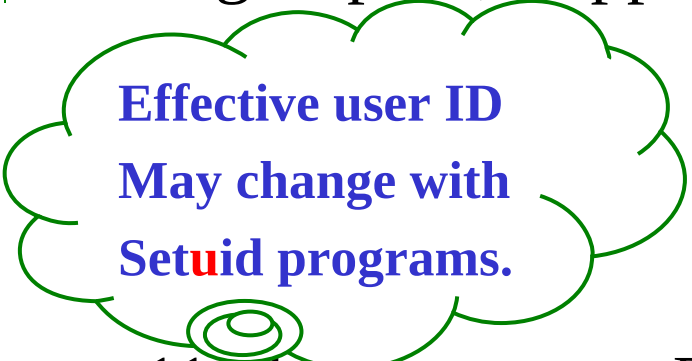


Zombies Revisited.




Running executables

- Typically an executable is loaded into the current process's address space.
- As a result the current process will start running the code from the executable.
- The old code of the program that loaded the executable will no longer execute.
- Loading executable will not result in any change to PID, PPID, Real user ID, Real group Id, Process group ID, Supplementary group Ids.



Effective user ID
May change with
Setuid programs.



Effective group ID
May change with
Setgid programs.

Logic Option

Exec family of library calls

- The library calls **execl()**, **execlp()**, **execle()**
execv(), **execvp()**

What about
execve() ?

Constitute the exec family of library calls.

- The exec class of calls replace the calling process's address space with a new executable file.

```
main()
```

```
{
```

```
printf("Before execl()\n");
```

```
execl ("/home/kumar/myprog", "myprog", (char *) 0);
```

```
/* The above calls loads the executable /home/kumar/myprog  
into the current process and starts executing it's main().
```

```
Exec functions never return on successful execution. */
```

```
printf("This printf() should not be executed.\n");
```

```
}
```

Execl

- `int execl(const char *path, const char *arg, ...);`
`execl ("/home/kumar/myprog",`
`"myprog", "arg1", "arg2", "arg3", (char *) 0);`
- `path` - pathname of the file to be executed in the current process.
- The program `myprog` will start executing in the current process. No change in PID or PPID.
- Assuming c Source for `myprog` , execution starts in `main()` of `myprog` . The command line arguments would be :

`"myprog", "arg1", "arg2", "arg3", (char *) 0`
`argv[0] argv[1] argv[2] argv[3] argv[4]`

Conventionally, Program name.
Any string allowed.

Execlp

- `int execlp(const char *file, const char *arg, ...);`
`execlp ("ls",
 "ls", "-l", "/etc/passwd", (char *) 0);`
- `file` - filename of the file to be executed in the current process.
- The Shell variable PATH will be used to locate a file with name "ls" that has execute permission.
- If found, that executable will start executing in the current process. No change in PID or PPID.
- The command line arguments would be :

"ls",	"-l",	"/etc/passwd",	(char *) 0
argv[0]	argv[1]	argv[2]	argv[3]

Execv

- `int execv(const char *path, char *const argv[]);`
`char * a[]={ "ls", "-l", "/etc/passwd", (char *) 0};`

`execv ("/bin/ls",a);`
- `path` – Path of filename of the file to be executed in the current process.
- Uses a vector of arguments instead of a list.
- Last element of the vector should be a null pointer.

<code>"ls",</code>	<code>"-l",</code>	<code>"etc/passwd",</code>	<code>(char *) 0</code>
<code>argv[0]</code>	<code>argv[1]</code>	<code>argv[2]</code>	<code>argv[3]</code>

Execvp

- `int execvp(const char *file, char *const argv[]);`
`char * a[]={ "ls", "-l", "/etc/passwd", (char *) 0};`

`execvp ("ls",a);`
- `file` – Filename of the file to be executed in the current process.
- Uses a vector of arguments instead of a list.
- Last element of the vector should be a null pointer.

<code>"ls",</code>	<code>"-l",</code>	<code>"/etc/passwd",</code>	<code>(char *) 0</code>
<code>argv[0]</code>	<code>argv[1]</code>	<code>argv[2]</code>	<code>argv[3]</code>

Combining fork and exec

```
if( (pid = fork ()) >0)
{
    wait( &status); /* Parent */
    if (WIFEXITED (status))
        printf ("Child terminated Normally.
}
else
{exec1 (  "/bin/ls",      /* child */
        "ls", "-l", "/etc/passwd", (char *) 0);

printf ("IF THIS IS PRINTED, \
        THERE WAS AN ERROR\n");
}
```

Redirecting standard devices..

```
if( (pid = fork ()) >0)
{ wait( &status); /* Parent */ }
else
{
    int fd = open("f1",
                  O_CREAT|O_WRONLY|O_TRUNC);
    close(1);
    dup(fd);
    execl (  "/bin/ls",      /* child */
           "ls", "-l", "/etc/passwd", (char *) 0);
}
```